Evaluating LLM Reliability for Coding Adversarial Machine Learning Research Papers

***Abstract*—Meta-research on security and machine learning often requires systematic coding of large paper corpora, a task that has remained manual and slow. Large language models (LLMs) may help automate annotation, but their reliability for coding security papers is unknown. We evaluate LLM reliability for coding adversarial ML research by comparing four models (GPT-4o, GPT-5.2, Gemini-3, Claude Sonnet-4.5) to human expert coding on 71 papers and 9 categorical variables. All tested models reach nearly perfect intra-rater consistency (Fleiss' $\kappa = 0.97$–1.00), high inter-LLM agreement ($\kappa = 0.79$–0.85), and substantial agreement with the human coder ($\kappa = 0.64$–0.71). Reliability varies by category: structured variables like domain reach almost perfect agreement ($\kappa = 0.86$), while interpretive ones like attack category are moderate ($\kappa = 0.55$). LLMs are usable as screening tools for structured meta-research in security, though interpretive judgment still benefits from human review.**

1. Introduction

The security and privacy research community increasingly recognizes the need for critical self-reflection on research practices [1], [2], [3]. From evaluating peer review processes [4] to assessing reproducibility [5], meta-research enables evidence-based improvements to how we conduct, evaluate, and sustain scientific work. Careful meta-analysis still requires someone to code papers: methodology, evaluation practices, artifacts, threat models. That does not scale; manual coding is costly and usually keeps studies to tens of papers.

LLMs have shown decent reliability on text annotation in social science [6], [7] and medical reviews [8], with some agreement levels close to human annotators [9]. Nobody has yet tested whether they can reliably code *security* papers, where threat models, evaluation methodology, and jargon are specific. Reliable LLM coding could support large-scale meta-analysis; if not, knowing where they fail clarifies when expert coding is still needed.

Adversarial ML is a good testbed. The field has known evaluation issues [10], [11], [12], [13] (overstated robustness, weak adaptive evaluation), critiqued in manual work [14] but never in a large systematic coding study. Reliable automated coding would make that kind of meta-analysis feasible.

Contribution. We run the first evaluation of LLM reliability for coding security papers and ask: (1) How stable is each LLM across repeated codings (intra-rater)? (2) How well do LLMs match human coders? (3) Which variables are reliable or brittle under automation?

We use four LLMs (GPT-4o, GPT-5.2, Google Gemini-3, Claude Sonnet-4.5) on 71 adversarial ML papers and 9 categorical variables: research type, domain, venue, artifacts, threat model, and evaluation. Each LLM coded all papers 3 times; we compute Fleiss' κ for intra-rater and Cohen's κ for inter-rater agreement [15], [16].

Results are encouraging. All four models achieve nearly perfect self-consistency ($\kappa = 0.97$–1.00), high inter-LLM agreement ($\kappa = 0.79$–0.85), and substantial agreement with the human coder ($\kappa = 0.64$–0.71), in the range often deemed acceptable [17]. By category, structured variables (e.g., domain $\kappa = 0.86$, threat model $\kappa = 0.75$) reach high agreement, while interpretive ones (e.g., attack category $\kappa = 0.55$) are moderate.

LLMs look useful as *screening* tools for structured variables and could cut human effort on a first pass; they still fall short of expert coding on hard methodological calls. We report reliability numbers and release codebooks, prompts, and scripts for reuse.

2. Related Work

Recent work demonstrates that LLMs can match or approach human annotator performance on text classification tasks in social science [6], [7], medical systematic reviews [8], and content moderation [9]. Gilardi et al. [6] find GPT-3.5 achieves human-level agreement ($\kappa \approx 0.7$) on tweet classification, with GPT-4 outperforming crowd workers. Ziems et al. [18] systematically compare LLMs to human annotators across 27 datasets, finding that LLMs excel on structured tasks but struggle with subjective judgments. Törnberg [7] had GPT-4 beat experts and crowd workers on political tweets, while Rathje et al. [9] got human-like correlations ($r = 0.59$–0.77) across 15 datasets and 12 languages ($n = 47{,}925$). For systematic reviews the picture is mixed: Khraisha et al. [19] got adjusted $\kappa = 0.65$ for GPT-4 on screening, while a scoping review of 37 studies [20] called LLMs "not yet ready" as sole review tools. Gartlehner et al. [8], on 9,341 extractions from 63 medical studies, found Claude-assisted extraction at 91% accuracy vs. 89% human-only. Known issues include model drift [21], prompt sensitivity [22], and extraction hallucinations. No prior work has evaluated LLM coding of *security* papers, where threat

models, evaluation methodology, and jargon are domain-specific.

Content analysis uses standard agreement metrics: Cohen's κ [23] handles two raters with chance correction, while Fleiss' κ [24] extends to multiple raters. Krippendorff's α [15] is flexible for any number of coders; common advice is $\alpha \geq 0.80$ for reliability [25]. Landis and Koch [17] label $\kappa = 0.61$–0.80 "substantial" and 0.81–1.0 "almost perfect." Neuendorf [16] suggests $\kappa \geq 0.80$ for published work. We use Fleiss' κ for intra-rater reliability and Cohen's κ for inter-rater, reporting bootstrap 95% confidence intervals. For threat-model variables (T1/T2), we compute subset κ excluding N/A to avoid prevalence inflation [26].

Adversarial ML has well-known evaluation problems. Carlini and Wagner [10] showed defensive distillation gave almost no robustness despite 99.5% claims. Athalye et al. [11] showed 7 of 9 non-certified ICLR 2018 defenses relied on "obfuscated gradients"; Tramèr et al. [12] broke 13 more from top venues. Arp et al. [13] reviewed 30 papers over 10 years and listed 10 pitfalls (data snooping, sampling bias, lab-only evaluation). These critiques say *what* is wrong but not *how common* each pattern is. Large-scale coding could answer that. Prior meta-research in security has examined peer review [4], reproducibility [5], and artifact evaluation [3], but typically analyzes only tens of papers due to manual coding costs. Olszewski et al. [5] found no clear gain in artifact availability after artifact-evaluation committees appeared (around 2020) across roughly 750 ML security papers. Reliable LLM coding could enable systematic reviews of hundreds of papers to quantify trends across entire subfields.

3. Methodology

3.1. Paper Selection

We selected 71 adversarial ML papers based on adoption by widely-used security artifacts as a proxy for research impact and visibility. Papers were included if they were cited, implemented, or referenced in major adversarial ML tools (CleverHans, IBM ART, Foolbox, TextAttack, PyRIT), benchmarks (RobustBench, AutoAttack, HarmBench), or industry threat frameworks (MITRE ATLAS). This artifact-adoption criterion ensures papers have demonstrable influence on practice, though it may overrepresent work with clearer reporting and available code relative to a random sample. Papers span 2013–2024 and cover foundational attacks (e.g., FGSM [27], C&W [28]), certified defenses [29], robustness benchmarks [30], and recent LLM jailbreaking work [31]. An arXiv ID and title identify each paper.

3.2. Coding and Data Extraction

We extract full text from each PDF with PyMuPDF (first 15 pages, max 50,000 characters). All raters saw the papers in the same order. For code availability variables (G5/G6), this PDF-only input restriction may disadvantage LLMs relative to the human coder, who could access external resources (GitHub, project pages).

Each paper was coded on 9 categorical variables grouped into three themes: research characteristics (G1–G6), threat protection properties specific to attack papers (T1–T2) and practical evaluation (Q1). These coding variables are standard and are commonly seen in the literature in adversarial ML reviews like [13], [32]; our framework was inspired in particular by [32]. Table 1 includes the complete variable list and allowed values. Some categories intentionally overlap (e.g., "Attack/Defense") to capture papers spanning multiple contributions, reflecting real-world ambiguity in the literature.

We use four LLMs and one human rater (Table 2). Based on their state-of-the-art performance in coding papers, we chose GPT-4o, GPT-5.2, Gemini-3 and Claude Sonnet-4.5 for coding our 71 papers. Each LLM was prompted separately through independent API calls with the same codebook, definition, allowed values, and a 0 temperature for consistency. We made 3 runs for each of the four models. A single human rater coded all 71 papers based on the same codebook.

Before computing agreement, we normalised all the values by trimming whitespace, hyphens, underscores, and removing `None`, `N/A`, `null`, and empty values, collapsing them to `n/a` so that formatting noise does not cause disagreement. We use two standard metrics [15], [16] to compute agreement. For models with 3 runs, we compute Fleiss' κ [24] over the three runs with items = paper $\times$ variable. This measures the model self-consistency across different runs. For the Human coder with only 1 run, intra-rater reliability cannot be computed and is marked as N/A. For each pair of raters, we compute Cohen's κ [23] on coding consensus. We concatenate all 9 columns into 639 labels per rater (71×9) for a global 5×5 matrix (diagonal = Fleiss' κ for multi-run models, off-diagonal = Cohen's κ). We also report mean κ per variable (G1–G6, T1–T2, Q1) across pairs to see which codes are stable. Interpretation follows the Landis and Koch [17] scale: $\kappa < 0.21$ (slight), 0.21–0.40 (fair), 0.41–0.60 (moderate), 0.61–0.80 (substantial), 0.81–1.00 (almost perfect).

4. Results

Figure 1 shows the full agreement matrix. Diagonal entries report intra-rater reliability (Fleiss' κ); off-diagonal entries report inter-rater agreement (Cohen's κ). Each κ is computed over all 639 labels (71 papers $\times$ 9 variables) for the given rater pair.

All four LLMs show nearly perfect self-consistency (intra-rater reliability): Gemini-3 achieves $\kappa = 1.000$ (identical outputs across 3 runs), Sonnet-4.5 reaches $\kappa = 0.994$, while GPT-4o and GPT-5.2 both reach $\kappa = 0.975$. Temperature-0 sampling produces highly deterministic behavior. For inter-rater reliability, all four LLMs show high inter-model agreement ($\kappa = 0.79$–0.85), suggesting they converge on similar interpretations. The best LLM pair is GPT-5.2 $\leftrightarrow$ Gemini-3 ($\kappa = 0.846$), while Sonnet-4.5

TABLE 1. CODING SCHEMA: 9 CATEGORICAL VARIABLES

Variable	Description	Allowed Values
Group 1: Research Characteristics		
G1	Paper type	Attack, Defense, Evaluation, Both, Attack/Defense, Attack/Evaluation, Defense/Evaluation, Attack/Defense/Evaluation
G2	Attack category (if attack)	Evasion, Poisoning, Privacy, Defense, N/A
G3	Application domain	Vision, NLP, LLM, Audio, Malware, Tabular, Cross, Cross-domain
G4	Publication venue category	ML, Security, Journal, arXiv-only
G5	Code currently available	Yes, No
G6	Code release timing	At-pub, Post-pub, Never
Group 2: Threat Model (Attack Papers Only)		
T1	Adversary model access	White, Black, Gray, White/Black, N/A
T2	Gradient computation required	Yes, No, N/A
Group 3: Practical Evaluation		
Q1	Evaluation on real-world system	Yes, No, Partial, N/A

TABLE 2. RATERS AND RUN CONFIGURATION

Rater	Model/Source	Runs	Consensus
GPT-4o	OpenAI API	3	Majority vote
GPT-5.2	OpenAI API	3	Majority vote
Gemini-3	Google Gemini 3.0 Pro	3	Majority vote
Sonnet-4.5	claude-sonnet-4-5-20250929	3	Majority vote
Human	Domain expert	1	N/A

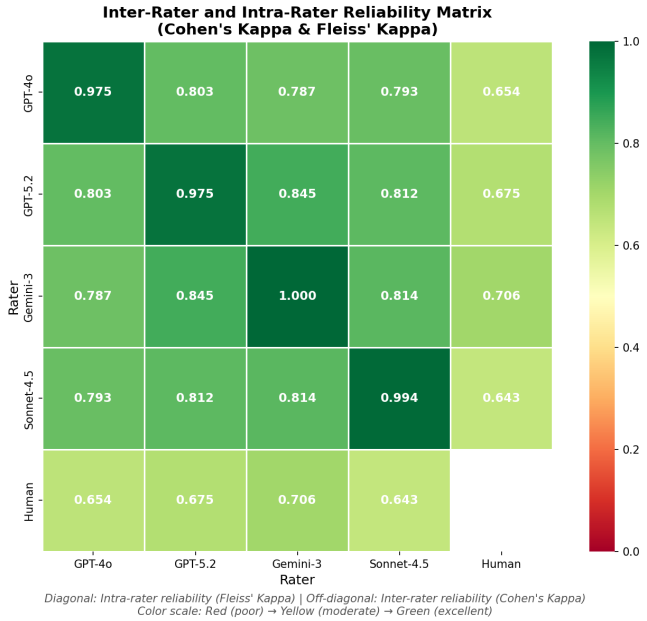

Figure 1. Global agreement matrix. Diagonal (dark green): intra-rater Fleiss' κ for all LLMs (3 runs each); N/A for the single-run Human coder. Off-diagonal: inter-rater Cohen's κ. Color scale: red (poor) $\rightarrow$ green (excellent).

achieves excellent agreement with other LLMs ($\kappa = 0.79$–0.81). All LLMs achieve substantial agreement with the human coder: Gemini-3 ($\kappa = 0.707$), GPT-5.2 ($\kappa = 0.675$), GPT-4o ($\kappa = 0.654$), and Sonnet-4.5 ($\kappa = 0.643$), all in the "substantial" range that Landis and Koch [17] consider acceptable for published work. Table 3 shows mean κ across all 10 inter-rater pairs, ranked by reliability; per-variable heatmaps appear in Appendix A.

TABLE 3. PER-VARIABLE AGREEMENT (MEAN κ ACROSS ALL PAIRS WITH 95% CI)

Var	Description	Mean κ	95% CI	Interp.
G3	Application domain	.856	[.802, .901]	Almost perfect
T1	Threat model access	.748	[.674, .819]	Substantial
G5	Code available	.742	[.658, .821]	Substantial
T2	Gradient required	.736	[.644, .822]	Substantial
G6	Code release timing	.699	[.614, .780]	Substantial
Q1	Real-world evaluation	.614	[.508, .716]	Moderate
G4	Publication venue	.611	[.482, .734]	Moderate
G1	Paper type	.589	[.469, .706]	Moderate
G2	Attack category	.548	[.418, .674]	Moderate

Structured variables perform best: application domain (G3) achieves almost perfect agreement ($\kappa = 0.86$), while threat-model access (T1), gradient requirement (T2), code availability (G5), and release timing (G6) all reach substantial agreement ($\kappa = 0.70$–0.75). These have well-defined categories with clear textual signals (e.g., "white-box," "black-box," "vision," "NLP"). Interpretive variables are harder: publication venue (G4), real-world evaluation (Q1), paper type (G1), and attack category (G2) achieve moderate agreement ($\kappa = 0.55$–0.61). These require judgment. For example, is a paper that proposes an attack and then defends against it "Attack," "Defense," or "Both"? Is a poisoning attack also "privacy" if it extracts training data? Disagreements arise from ambiguous boundaries. The best pairs per variable reach substantial to almost perfect agreement: G3 (domain) achieves $\kappa = 0.87$ (GPT-5.2 $\leftrightarrow$ Sonnet-4.5), T2 (gradient) reaches $\kappa = 0.90$ (GPT-5.2 $\leftrightarrow$ Gemini-3), and G5 (code available) hits $\kappa = 0.89$ (GPT-5.2 $\leftrightarrow$ Gemini-3). Even G2 (attack category) achieves $\kappa = 0.79$ between GPT-5.2 and Sonnet-4.5.

One caveat concerns threat-model variables: while T1/T2 appear to reach substantial agreement in Table 3,

subset κ excluding N/A reveals much lower LLM–human agreement ($\kappa = 0.18$–0.33). LLMs marked 56–70% of papers as N/A for T1/T2, while the human marked only 41–45%, indicating these variables are less reliable than global metrics suggest and benefit from N/A-driven prevalence inflation (see Limitations).

5. Discussion

5.1. Why Application Domain Works and Paper Type Fails

The core finding is not just that some variables reach high agreement and others do not, but that the gap reveals what LLMs can and cannot extract from text. Application domain (G3, $\kappa = 0.86$) works because it maps to concrete, repeated lexical signals. Papers testing on ImageNet, CIFAR-10, or MNIST reliably mention these datasets in abstracts and evaluation sections. Vision papers contain terms like "image," "CNN," and "pixel"; NLP papers mention "text," "embedding," or specific models like BERT. These are bag-of-words features that correlate strongly with ground truth, and LLMs trained on millions of papers have learned these associations. The task is pattern matching, not interpretation.

Paper type (G1, $\kappa = 0.59$) and attack category (G2, $\kappa = 0.55$) fail for a different reason. They require resolving ambiguity at the conceptual level, not the word level. A paper presenting AutoAttack is simultaneously an attack method (it generates adversarial examples) and an evaluation framework (it benchmarks defenses). Whether the human coder labels it "Attack" or "Evaluation" depends on what they infer the primary contribution to be, which requires reading the abstract, understanding author framing, and making a judgment call. LLMs see the surface form but struggle with the metalevel question of authorial intent. The codebook allows overlapping categories ("Attack/Evaluation"), but this only shifts the ambiguity; it does not resolve it.

Real-world evaluation (Q1, $\kappa = 0.61$) sits in the middle. LLMs can detect explicit mentions of deployed systems ("Google Cloud Vision API," "Tesla Autopilot," "GPT-4") with high accuracy, which is lexical matching. But they struggle with the boundary between "Partial" and "No." A paper testing on a realistic traffic-sign dataset in a lab setting is not deployed, but it is closer to real-world conditions than CIFAR-10. The human coder applies a gradient of realism that depends on background knowledge about how adversarial ML transfers from benchmarks to deployment. LLMs lack that contextual model and revert to surface cues, sometimes misclassifying realistic simulations as "Yes" if the paper emphasizes real-world relevance in the abstract.

5.2. The Threat Model Paradox

Threat model variables (T1, T2) appear to perform well ($\kappa = 0.75$, $\kappa = 0.74$), but this is misleading. When we compute κ excluding N/A, agreement drops to $\kappa = 0.18$–0.33 for LLM-human pairs. The high global κ is driven by agreement on which papers are *not* attacks (hence N/A), not by reliable extraction of threat model properties from attack papers. LLMs over-assign N/A (56–70% of papers) compared to the human (41–45%). This reflects a conservative strategy: when uncertain whether a paper is an attack, mark T1/T2 as N/A. This inflates agreement through what Feinstein and Cicchetti [26] call "prevalence-induced" κ, where most items fall into one category (N/A), making chance-corrected agreement artificially high. The lesson is that LLMs can identify papers that clearly describe gradient-based white-box attacks (the textbook case), but struggle with boundary cases like defenses that include ablation attacks, or evaluation papers that propose new attacks as baselines.

5.3. What This Means for Automated Meta-Research

The pattern is clear: LLMs reliably automate extraction tasks that reduce to lexical pattern matching over well-scoped categories. Application domain, code availability ("GitHub," "code"), and gradient use ("∇," "backpropagation") are extractable because they have stable, repeated surface forms. Coding tasks that require resolving conceptual ambiguity, applying unstated background knowledge, or inferring author intent remain unreliable. This matches findings from social science annotation [18], where sentiment and stance detection (lexical) outperform intent classification (interpretive).

For security meta-research, this suggests a practical division of labor. Use LLMs to pre-code structured metadata (venue, domain, artifacts, threat model access) at scale, cutting human effort by 70–80%. Reserve human coding for interpretive variables (paper type, primary contribution, evaluation rigor) where disagreement reflects genuine ambiguity in the literature. Our results ($\kappa = 0.64$–0.71 with humans) indicate LLM output is a useful first pass, not ground truth. The right workflow is LLM-assisted coding with human adjudication, not full automation.

5.4. Implications for Adversarial ML Evaluation Studies

These findings carry broader implications for adversarial ML meta-research. Our 71-paper corpus serves as a proof of concept demonstrating that reliable LLM coding could enable systematic studies of evaluation practices, artifact availability, and threat-model reporting across hundreds of papers, the kind of quantitative meta-analysis that manual coding has historically precluded. Researchers could, for instance, track whether artifact-evaluation policies [33] correlate with actual code availability over time, or examine whether vulnerabilities such as "obfuscated gradients" [11] persist in recent defenses. Our reliability numbers provide the first evidence that such large-scale studies are feasible for structured variables in security research.

6. Limitations

One of the major limitations of our work arises from resource constraints. First, we only use one expert coder due to the large number of papers and coding criteria. A single human coder may have bias, which can be mitigated by averaging or majority voting among three human raters. Prior work suggests human–human κ for similar tasks is 0.6–0.8 [15]; our LLM–human $\kappa = 0.71$ is strong, but this should be validated with multiple humans. Second, we were able to test only a few state-of-the-art models due to budget limitations for API calls. We would need to test in future with other models which show promise: Opus 4.6, o3 reasoning. As these models keep evolving, future models might show different results than models in early 2026. Third, we only used a single one-shot prompt and did not optimise. Per-model tuning or a few-shot examples could improve reliability. Fourth, the intra-rater consistency of models across three runs is not particularly significant and is largely a consequence of keeping the temperature to 0. Finally, the models might be influenced by their training data on the papers we included. The choice of top adversarial ML papers may not generalise to other security fields (systems security, cryptography).

7. Conclusion

We evaluated LLM reliability for coding adversarial ML papers on 9 categorical variables and found that all models achieve nearly perfect self-consistency ($\kappa = 0.97$–1.00) and substantial agreement with a human expert ($\kappa = 0.64$–0.71), with reliability varying by variable type: structured codes like domain reach almost perfect agreement ($\kappa = 0.86$), while interpretive codes like attack category are moderate ($\kappa = 0.55$–0.61). These results indicate that LLMs are usable as screening tools for structured meta-research in security, though they do not yet replace expert coding for nuanced methodological judgment. All supplementary materials including codebook, prompts, analysis scripts, raw CSV outputs, and per-model coding results are available at: https://github.com/actionproject-madhav/llm-paper-coding-reliability.

8. Ethics Considerations

We evaluate LLM coding of public academic papers. No human subjects; all papers are published and openly or institutionally accessible. No privacy, consent, or harm issues. The goal is to improve how the security community does meta-research.

9. LLM Usage Considerations

LLMs were both *evaluated* (as coders) and *used* (literature search, running experiments, analysis). We treat their coding output as data and validate it against human coding. LLMs helped with literature search, LaTeX formatting, running experiments, and debugging. Some experiment code is AI-generated with human verification. We release prompts, scripts, and raw outputs for replication.

References

[1] "MetaCRiSP 2026: 1st Workshop on Metascience and Critical Reflections in Security & Privacy," https://metacrisp.org/, 2026, co-located with IEEE S&P 2026.

[2] C. Rossow, C. J. Dietrich, C. Grier, C. Kreibich, V. Paxson, N. Pohlmann, H. Bos, and M. van Steen, "Prudent practices for designing malware experiments: Status quo and outlook," in *Proceedings of the IEEE Symposium on Security and Privacy (S&P)*, 2012, pp. 65–79.

[3] A. Baset and T. Denning, "A data-driven reflection on 36 years of security and privacy research," in *Proceedings of the 12th USENIX Workshop on Cyber Security Experimentation and Test (CSET)*, 2019.

[4] V. Le Pochat, "Reflecting on research practices: Encouraging the security research community to engage in meta-research," *Communications of the ACM*, 2024, opinion, Computing Profession. Posted Mar. 27 2024.

[5] D. Olszewski, A. Lu, C. Stillman, K. Warren, C. Kitroser, A. Pascual, D. Ukirde, K. R. B. Butler, and P. Traynor, ""get in researchers; we're measuring reproducibility": A reproducibility study of machine learning papers in tier 1 security conferences," in *Proceedings of the ACM SIGSAC Conference on Computer and Communications Security (CCS)*, 2023, pp. 2349–2363.

[6] F. Gilardi, M. Alizadeh, and M. Kubli, "ChatGPT outperforms crowd workers for text-annotation tasks," *Proceedings of the National Academy of Sciences*, vol. 120, no. 30, p. e2305016120, 2023.

[7] P. Törnberg, "ChatGPT-4 outperforms experts and crowd workers in annotating political twitter messages with zero-shot learning," *arXiv preprint arXiv:2304.06588*, 2024.

[8] G. Gartlehner, S. Kugley, K. Crotty *et al.*, "Artificial intelligence–assisted data extraction with a large language model: A study within reviews," *Annals of Internal Medicine*, vol. 178, pp. 1763–1771, 2025, epub 4 November 2025. doi: 10.7326/ANNALS-25-00739.

[9] S. Rathje, D.-M. Mirea, I. Sucholutsky, R. Marjieh, C. Robertson, and J. J. Van Bavel, "GPT is an effective tool for multilingual psychological text analysis," *Proceedings of the National Academy of Sciences*, vol. 121, no. 8, p. e2308950121, 2024.

[10] N. Carlini and D. Wagner, "Towards evaluating the robustness of neural networks," in *Proceedings of the IEEE Symposium on Security and Privacy (S&P)*, 2017, pp. 39–57.

[11] A. Athalye, N. Carlini, and D. Wagner, "Obfuscated gradients give a false sense of security: Circumventing defenses to adversarial examples," in *Proceedings of the International Conference on Machine Learning (ICML)*, 2018, pp. 274–283.

[12] F. Tramèr, N. Carlini, W. Brendel, and A. Madry, "On adaptive attacks to adversarial example defenses," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, 2020, pp. 1633–1645.

[13] D. Arp, E. Quiring, C. Wressnegger, and K. Rieck, "Dos and don'ts of machine learning in computer security," in *Proceedings of the USENIX Security Symposium*, 2022, pp. 3971–3988, distinguished Paper Award.

[14] N. Carlini, A. Athalye, N. Papernot, W. Brendel, J. Rauber, D. Tsipras, I. Goodfellow, A. Madry, and A. Kurakin, "On evaluating adversarial robustness," *arXiv preprint arXiv:1902.06705*, 2019.

[15] K. Krippendorff, *Content Analysis: An Introduction to Its Methodology*, 2nd ed. Sage Publications, 2004.

[16] K. A. Neuendorf, *The Content Analysis Guidebook*, 2nd ed. Sage Publications, 2017.

[17] J. R. Landis and G. G. Koch, "The measurement of observer agreement for categorical data," *Biometrics*, vol. 33, no. 1, pp. 159–174, 1977.

[18] C. Ziems, W. Held, O. Shaikh, J. Chen, Z. Zhang, and D. Yang, "Can large language models transform computational social science?" *Computational Linguistics*, vol. 50, no. 1, pp. 237–291, 2024.

[19] Q. Khraisha, S. Put, J. Kappenberg, A. Warraitch, and K. Hadfield, "Can large language models replace humans in systematic reviews? evaluating GPT-4's efficacy in screening and extracting data from peer-reviewed and grey literature in multiple languages," *Research Synthesis Methods*, 2024, early view. doi: 10.1002/jrsm.1715.

[20] J.-L. Lieberum, M. Toews, M.-I. Metzendorf, F. Heilmeyer, W. Siemens, C. Haverkamp, D. Böhringer, J. J. Meerpohl, and A. Eisele-Metzger, "Large language models for conducting systematic reviews: On the rise, but not yet ready for use—a scoping review," *Journal of Clinical Epidemiology*, vol. 181, p. 111746, 2025.

[21] L. Chen, M. Zaharia, and J. Zou, "How is ChatGPT's behavior changing over time?" *arXiv preprint arXiv:2307.09009*, 2023.

[22] M. Sclar, Y. Choi, Y. Tsvetkov, and A. Suhr, "Quantifying language models' sensitivity to spurious features in prompt design or: How I learned to start worrying about prompt formatting," *arXiv preprint arXiv:2310.11324*, 2023.

[23] J. Cohen, "A coefficient of agreement for nominal scales," *Educational and Psychological Measurement*, vol. 20, no. 1, pp. 37–46, 1960.

[24] J. L. Fleiss, "Measuring nominal scale agreement among many raters," *Psychological Bulletin*, vol. 76, no. 5, pp. 378–382, 1971.

[25] A. F. Hayes and K. Krippendorff, "Answering the call for a standard reliability measure for coding data," *Communication Methods and Measures*, vol. 1, no. 1, pp. 77–89, 2007.

[26] A. R. Feinstein and D. V. Cicchetti, "High agreement but low kappa: I. the problems of two paradoxes," *Journal of Clinical Epidemiology*, vol. 43, no. 6, pp. 543–549, 1990.

[27] I. J. Goodfellow, J. Shlens, and C. Szegedy, "Explaining and harnessing adversarial examples," in *Proceedings of the International Conference on Learning Representations (ICLR)*, 2015, arXiv:1412.6572.

[28] N. Carlini and D. Wagner, "Towards evaluating the robustness of neural networks," in *Proceedings of the IEEE Symposium on Security and Privacy (S&P)*, 2017, pp. 39–57.

[29] J. M. Cohen, E. Rosenfeld, and J. Z. Kolter, "Certified adversarial robustness via randomized smoothing," in *Proceedings of the International Conference on Machine Learning (ICML)*, 2019, pp. 1310–1320, arXiv:1902.02918.

[30] F. Croce and M. Hein, "Reliable evaluation of adversarial robustness with an ensemble of diverse parameter-free attacks," in *Proceedings of the International Conference on Machine Learning (ICML)*, 2020, arXiv:2003.01690.

[31] A. Zou, Z. Wang, N. Carlini, M. Nasr, J. Z. Kolter, and M. Fredrikson, "Universal and transferable adversarial attacks on aligned language models," *arXiv preprint arXiv:2307.15043*, 2023.

[32] G. Apruzzese, H. S. Anderson, S. Dambra, D. Freeman, F. Pierazzi, and K. A. Roundy, ""real attackers don't compute gradients": Bridging the gap between adversarial ML research and practice," in *Proceedings of the IEEE Conference on Secure and Trustworthy Machine Learning (SaTML)*, 2023, pp. 339–364.

[33] "USENIX Security '25 Call for Artifacts," https://www.usenix.org/conference/usenixsecurity25/call-for-artifacts, 2025, accessed: 2026-02-12.

Appendix

Appendix: Supplementary Material and Coding Prompt

Figure 2 shows Cohen's κ for each variable across all rater pairs. Diagonal entries show intra-rater reliability (Fleiss' κ) for all four LLMs (each with 3 runs); the Human coder diagonal is N/A (single run). Variables G3 (domain), G5 (code available), T1 (threat model), and T2 (gradient) show consistent high agreement across most LLM pairs. Variables G1 (paper type), G2 (attack category), and Q1 (real-world evaluation) show moderate agreement, reflecting their interpretive nature.

All LLMs received identical system prompts with temperature set to 0.0 for deterministic outputs. Each paper was provided as extracted PDF text (first 15 pages, max 50,000 characters) with title and arXiv ID. The 71 papers were selected based on artifact-adoption criteria detailed in Section 3.1.

All supplementary materials are available at our GitHub repository: https://github.com/actionproject-madhav/llm-paper-coding-reliability, including raw CSV outputs from each model run (3 runs per LLM), human expert codings, the complete paper list with arXiv IDs and PDF links, analysis scripts for computing κ statistics, and API interaction code for each model.

Coding Prompt

```
You are coding academic papers for adversarial ML research. For each paper, assign values to exactly 9 variables. Use
     ONLY the allowed values listed below.

## CODING VARIABLES (STRICT VALUES)

### G1: Paper Type
ALLOWED: "Attack", "Defense", "Evaluation", "Both", "Attack/Defense", "Attack/Evaluation", "Defense/Evaluation", "Attack
     /Defense/Evaluation"

### G2: Threat Category
ALLOWED: "Evasion", "Poisoning", "Privacy", "Defense", "N/A"

### G3: Domain
ALLOWED: "Vision", "NLP", "LLM", "Audio", "Malware", "Tabular", "Cross", "Cross-domain"

### G4: Publication Venue
ALLOWED: "ML", "Security", "Journal", "arXiv-only"

### G5: Code Available
ALLOWED: "Yes", "No"

### G6: Code Release Timing
ALLOWED: "At-pub", "Post-pub", "Never"

### T1: Model Access Level (attacks only)
ALLOWED: "White", "Black", "Gray", "White/Black", "N/A"

### T2: Gradient Required (attacks only)
ALLOWED: "Yes", "No", "N/A"

### Q1: Real-World Evaluation
ALLOWED: "Yes", "No", "Partial", "N/A"

## OUTPUT FORMAT
Return ONLY valid JSON (no markdown):
{
  "G1": "<value>", "G2": "<value>",
  "G3": "<value>", "G4": "<value>",
  "G5": "<value>", "G6": "<value>",
  "T1": "<value>", "T2": "<value>",
  "Q1": "<value>",
  "reasoning": "Brief explanation"
}

CRITICAL: Use ONLY the allowed values. Do NOT create new values.
```

Figure 2. Per-variable agreement heatmaps showing pairwise κ for each coding variable (G1–G6: research characteristics; T1–T2: threat model; Q1: evaluation). Diagonal entries show intra-rater Fleiss' κ for all LLMs; Human is N/A (single run). Color scale ranges from red (poor) to green (excellent).